

```
In [1]: import pandas as pd
import numpy as np
from collections import defaultdict
import re
```

Cleaning function:

```
In [2]: def preprocess_string(str_arg):

    cleaned_str=re.sub('[^a-z\s]+',' ',str_arg,flags=re.IGNORECASE)
    cleaned_str=re.sub('\s+',' ',cleaned_str)
    cleaned_str=cleaned_str.lower()

    return cleaned_str
```

Naive Bayes from scratch:

```
In [3]: class NaiveBayes:

    def __init__(self,unique_classes):
        self.classes=unique_classes

    def addToBow(self,example,dict_index):

        if isinstance(example,np.ndarray): example=example[0]

        for token_word in example.split():
            self.bow_dicts[dict_index][token_word]+=1

    def fit(self,dataset,labels):
```

```

        self.examples=dataset
        self.labels=labels
        self.bow_dicts=np.array([defaultdict(lambda:0) for index in range(self.classes.shape[0])])

        for cat_index,cat in enumerate(self.classes):
            all_cat_examples=self.examples[self.labels==cat]
            cleaned_examples=[preprocess_string(cat_example) for cat_example in all_cat_examples]
            cleaned_examples=pd.DataFrame(data=cleaned_examples)
            np.apply_along_axis(self.addToBow,1,cleaned_examples,cat_index)

        prob_classes=np.empty(self.classes.shape[0])
        all_words=[]
        cat_word_counts=np.empty(self.classes.shape[0])

        for cat_index,cat in enumerate(self.classes):
            prob_classes[cat_index]=np.sum(self.labels==cat)/float(self.labels.shape[0])
            count=list(self.bow_dicts[cat_index].values())
            cat_word_counts[cat_index]=np.sum(np.array(list(self.bow_dicts[cat_index].values())))+1
            all_words+=self.bow_dicts[cat_index].keys()

        self.vocab=np.unique(np.array(all_words))
        self.vocab_length=self.vocab.shape[0]

        denoms=np.array([cat_word_counts[cat_index]+self.vocab_length+1 for cat_index,cat in enumerate(self.classes)])

        self.cats_info=[(self.bow_dicts[cat_index],prob_classes[cat_index],denoms[cat_index]) for cat_index,cat in enumerate(self.classes)]
        self.cats_info=np.array(self.cats_info)

```

```

def getExampleProb(self, test_example):

    likelihood_prob=np.zeros(self.classes.shape[0])
    for cat_index,cat in enumerate(self.classes):
        for test_token in test_example.split():

            test_token_counts=self.cats_info[cat_index][0].get(test
_token,0)+1
            test_token_prob=test_token_counts/float(self.cats_info[
cat_index][2])
            likelihood_prob[cat_index]+=np.log(test_token_prob)
        post_prob=np.empty(self.classes.shape[0])

        for cat_index,cat in enumerate(self.classes):
            post_prob[cat_index]=likelihood_prob[cat_index]+np.log(self
.cats_info[cat_index][1])
        return post_prob

def predict(self, test_set):

    predictions=[]
    for example in test_set:
        cleaned_example=preprocess_string(example)

        post_prob=self.getExampleProb(cleaned_example)
        predictions.append(self.classes[np.argmax(post_prob)])
    return np.array(predictions)

```

Training data-set:

```

In [4]: training_set=pd.read_csv ("train.csv")
training_set.info()
training_set.shape

```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60000 entries, 0 to 59999
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id          60000 non-null  int64
1   comment     60000 non-null  object
2   subreddit   60000 non-null  object
dtypes: int64(1), object(2)
memory usage: 937.6+ KB
```

Out[4]: (60000, 3)

In [5]: training_set.head()

Out[5]:

	id	comment	subreddit
0	0	I think prestige points should not expire ever...	leagueoflegends
1	1	Whats going to happen with them if they will b...	europe
2	2	Anecdotal evidence is anecdotal. Clearly by "e...	gameofthrones
3	3	Look dude, with all due respect, your music is...	Music
4	4	Hope he gets the doomhammer back!	wow

Prediction on train data-set:

```
In [6]: y_train=training_set['subreddit'].values
x_train=training_set['comment'].values

from sklearn.model_selection import train_test_split
train_data,test_data,train_labels,test_labels=train_test_split(x_train,
y_train,
                                                                    shuffle=
True,
```

```

e=0.2
state=1
y=y_train
classes=np.unique(train_labels)

nb=NaiveBayes(classes)
nb.fit(train_data,train_labels)

y_pred_train=nb.predict(test_data)
test_acc=np.sum(y_pred_train==test_labels)/float(test_labels.shape[0])

print ("Test Set Accuracy: ",test_acc)

```

Test Set Accuracy: 0.5183333333333333

Testing data-set:

```

In [7]: testing_set=pd.read_csv ("test.csv")
testing_set.info()
testing_set.shape

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    id         20000 non-null  int64
1    comment    20000 non-null  object
dtypes: int64(1), object(1)
memory usage: 234.4+ KB

```

Out[7]: (20000, 2)

```

In [8]: testing_set.head()

```

Out[8]:

	id	comment
0	0	Holy shit a shot counter.
1	1	It doesn't matter that it isn't hard to rememb...
2	2	I find it funny that this is downvoted
3	3	They are really getting ridiculous with all t...
4	4	He's Eden's best friend

Prediction on test data-set:

```
In [9]: X_test=testing_set.comment.values
        y_pred_test=nb.predict(X_test)
```

Submission:

```
In [10]: submission = zip(list(range(len(y_pred_test))), y_pred_test)
        test_df = pd.DataFrame(submission, columns=['Id', 'Category'])
        test_df.to_csv('submission.csv', index = False, header=True)
```

Naive Bayes cross validation by using KFold:

```
In [11]: from sklearn.model_selection import KFold

        X = x_train
        y = y_train
        kf = KFold(n_splits=5, random_state=42, shuffle=True)

        print(kf)
```

```
KFold(n_splits=5, random_state=42, shuffle=True)
```

```
In [12]: main_cross_val_accuracy = []
         for train_index, test_index in kf.split(X):
             X_train, X_test = X[train_index], X[test_index]
             y_train, y_test = y[train_index], y[test_index]
             nb.fit(X_train, y_train)
             y_pred_cross_train = nb.predict(X_test)
             test_cross_acc = np.sum(y_pred_cross_train == y_test) / float(y_test.shape[0])
             main_cross_val_accuracy.append(test_cross_acc)
```

```
In [13]: main_cross_val_accuracy
```

```
Out[13]: [0.4533333333333333,
          0.44683333333333336,
          0.45108333333333334,
          0.45466666666666666,
          0.44891666666666667]
```

```
#
```

Preparing the data for selective classifiers:

```
In [14]: import string
         from nltk.corpus import stopwords
         from nltk.tokenize import word_tokenize

         def clean_text(training_set):

             all_comments = list()
             lines = training_set["comment"].values.tolist()
             for text in lines:
                 text = text.lower()
```

```

pattern = re.compile('http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+]|[*\(\),]|(?:%[0-9a-fA-F][0-9a-fA-F]))+')
text = pattern.sub("", text)

text = re.sub(r"[.,\\"!@#$$%^&*(){}?/;`~:<>+=-]", "", text)

tokens = word_tokenize(text)

table = str.maketrans('', '', string.punctuation)

stripped = [w.translate(table) for w in tokens]
words = [word for word in stripped if word.isalpha()]

stop_words = set(stopwords.words("english"))
stop_words.discard("not")

words = [w for w in words if not w in stop_words]
words = ' '.join(words)

all_comments.append(words)
return all_comments

all_comments = clean_text(training_set)
all_comments[0:2]

```

Out[14]: ['think prestige points not expire ever skins buy available set duration example year release another skin vault old one making also limited edition skin also please love god not rerelease skins need grind prestige shop would suck everyone grinded',
'whats going happen refused asilum appeal']

Most frequent used words

```

In [15]: from nltk.probability import FreqDist
from nltk.stem import WordNetLemmatizer

```



```

c = all_comments
filtered_sentence = []
freq_count_limit = FreqDist()
lemmatizer=WordNetLemmatizer()
stop_words = set(stopwords.words('english'))

for i in c:
    comment_tokens = word_tokenize(i)

    for words in comment_tokens:
        if words not in stop_words:
            filtered_sentence.append(words)

            limit_words = lemmatizer.lemmatize(words)
#         for word in root_words:
            freq_count_limit[limit_words.lower()] += 1
freq_count_limit

```


Out[15]: FreqDist({'nt': 12955, 'like': 8953, 'one': 5971, 'people': 5884, 'get': 5839, 'would': 5745, 'think': 4498, 'time': 4385, 'game': 3900, 'even': 3839, ...})

In [16]: `import matplotlib.pyplot as plt`

```

freq_count_limit.plot(30,cumulative=False)
plt.show()

```



Vectorizing and transforming the text:

```
In [17]: from sklearn.feature_extraction.text import TfidfVectorizer

vect = TfidfVectorizer(ngram_range=(1,1), max_features=30000, strip_accents='ascii')
vect.fit(all_comments)
vocabulaire = vect.get_feature_names()
```

```
In [18]: bag_of_words = vect.transform(all_comments)
bag_of_words.shape
```

```
Out[18]: (60000, 30000)
```

Random Forest:

```
In [19]: from sklearn.ensemble import RandomForestClassifier
```

```
clf_rf = RandomForestClassifier(max_depth=15, random_state=42)
clf_rf.fit(bag_of_words, training_set['subreddit'])
```

Out[19]: RandomForestClassifier(max_depth=15, random_state=42)

```
In [20]: clf_rf.score(bag_of_words, training_set['subreddit'])
```

Out[20]: 0.34563333333333335

Random Forest - Cross Validation:

```
In [21]: from sklearn.model_selection import cross_val_score
```

```
scores_rf = cross_val_score(clf_rf, bag_of_words, training_set['subredd
it'], cv=5)
scores_rf
```

Out[21]: array([0.30675 , 0.30808333, 0.31591667, 0.31841667, 0.31941667])

Logistic Regression:

```
In [22]: from sklearn.linear_model import LogisticRegression
```

```
clf_lr = LogisticRegression(max_iter=10, random_state=42)
clf_lr.fit(bag_of_words, training_set['subreddit'])
```

c:\users\rezam\appdata\local\programs\python\python38-32\lib\site-packa
ges\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs fa
iled to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown
in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression  
n_iter_i = _check_optimize_result(
```

Out[22]: LogisticRegression(max_iter=10, random_state=42)

Logistic Regression - Cross Validation:

```
In [23]: clf_lr.score(bag_of_words, training_set['subreddit'])
```

Out[23]: 0.4988166666666667

Logistic Regression - Cross Validation:

```
In [25]: scores_lr = cross_val_score(clf_lr, bag_of_words, training_set['subreddit'], cv=5)  
scores_lr
```

```
c:\users\rezam\appdata\local\programs\python\python38-32\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression  
n_iter_i = _check_optimize_result(  
c:\users\rezam\appdata\local\programs\python\python38-32\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
c:\users\rezam\appdata\local\programs\python\python38-32\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
c:\users\rezam\appdata\local\programs\python\python38-32\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
c:\users\rezam\appdata\local\programs\python\python38-32\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
regression
n_iter_i = _check_optimize_result(
```

```
Out[25]: array([0.41966667, 0.41683333, 0.42291667, 0.42075    , 0.42783333])
```

```
In [ ]:
```